



33

Must-Ask Questions When Evaluating AI SOC Analysts



Coverage

1. Use cases

Which types of security alerts (cloud, identity, endpoint, email, network, data loss prevention) is the AI SOC Analyst able to triage, investigate, and remediate?



Coverage

2. Custom detections

Can the AI SOC Analyst investigate custom detections to address unique threats or organization-specific scenarios?



Coverage

3. New integrations

How quickly can the AI SOC Analyst support new alert types or security data sources to maintain consistent coverage and performance as the number of tools, alerts, and data sources grow?



Coverage

4. Threat hunting

Does the AI SOC Analyst support proactive threat hunting via natural language queries? Does this enable hypothesis-driven searches to see patterns and anomalies indicative of hidden threats or suspicious actions?



Accuracy

5. True or false positive

How accurate is the AI SOC Analyst in investigating alerts and identifying them as true positives and false positives? Are there any mechanisms in place to review false positives and measure their efficacy?



Accuracy

6. Reducing false positives

How effectively does the AI SOC Analyst reduce the number of false positives that a human analyst has to triage and investigate manually?



Accuracy

7. Hallucinations

What safeguards are in place to minimize the risk of hallucinations from LLMs?



Accuracy

8. Limited context

How does the AI SOC Analyst approach alerts when important contextual data is unavailable? What methods are used to prioritize, escalate, or defer these cases to maintain operational integrity?



Quality

9. Investigation depth

What questions does the AI SOC Analyst answer during an investigation, and how thoroughly does it emulate an expert human analyst to uncover every key detail?



Quality

10. Alert deduplication

How does the AI SOC Analyst identify and group duplicate alerts from the same event? Provide examples of how it reduces noise and streamlines investigations.



Quality

11. Incident view

How does the AI SOC Analyst group alerts from multiple security tools into an incident?



Quality

12. Follow-up questions

Does the AI SOC Analyst allow analysts to ask follow-up questions within an alert or across alerts using a natural language interface?



Quality

13. Transparency

How does the AI SOC Analyst outline the reasons behind its determination and severity of an alert? Are the steps and data points leading to a decision clearly outlined?



Quality

14. Measuring impact

What metrics are used to measure the effectiveness of the AI SOC Analyst and do they align with the customer's SOC priorities?



Quality

15. Enhanced detections

Can the AI SOC Analyst identify high-volume noisy alerts, gaps in coverage, or areas requiring tuning? Does it uncover insights to help security engineering teams improve detection strategies?



Quality

16. Response actions

How does the AI SOC Analyst enable faster resolution of incidents? What specific features or outputs reduce manual response efforts and improve response times?



Integrations

17. Workflow integrations

How does the AI SOC Analyst ensure seamless integration with tools like Slack or Microsoft Teams? Can it start workflows like notifying teams, escalating incidents, or invoking SOAR playbooks for remediation?



Integrations

18. Case management integrations

How does the AI SOC Analyst integrate with case management tools like Jira or ServiceNow? Does it allow analysts to create, update, and close tickets from the platform itself?



Integrations

19. API capabilities

Does the AI SOC Analyst offer APIs for customization and seamless integration with custom tools?



Customizability

20. Leveraging existing runbooks

Does the AI SOC Analyst provide the option to leverage existing runbooks for its investigations?



Customizability

21. Tailored investigations

Can customers tailor the AI SOC Analyst's investigation process, including defining steps, selecting data sources, and adjusting analysis, scoring, and impact assessments, to align with their specific needs?



Customizability

22. Granular feedback

Can users provide feedback at specific investigation steps or decision points?



Customizability

23. Global feedback

Can feedback be applied broadly across investigations to refine the AI SOC Analyst's analysis and decision-making?



Customizability

24. Impact of feedback

How quickly and effectively does the AI SOC Analyst incorporate user feedback to improve future investigations?



Time-to-value

25. Integration with tools and platforms

Which security tools and platforms does the AI SOC Analyst integrate with out-of-the-box, and how deep are the integrations?



Time-to-value

26. Deployment timeline

What is the average time required to deploy the AI SOC Analyst and begin processing alerts in a production environment?



Time-to-value

27. Initial results timeline

How quickly after deployment can the AI SOC Analyst provide measurable results in terms of reducing alert dwell time, improving mean time to investigate, or streamlining triage processes?



Time-to-value

28. Level of autonomy

How much assistance does a human analyst provide the AI SOC Analyst in triaging and investigating alerts?

Does it require pre-built playbooks?

Can it support both analyst-in-the-loop and autonomous investigations?



Time-to-value

29. Performance under testing/POV conditions

How well does the AI SOC Analyst perform during red team exercises, simulated attacks, or a proof-of-value (POV) assessment?



— Data privacy and security —

30. Use of customer LLMs

Does the AI SOC Analyst support the use of customer's LLM models?



— Data privacy and security —

31. Deployment options

Does the AI SOC Analyst offer the ability to run within a customer-managed environment? How can customers retain control of their data?



— Data privacy and security —

32. AI output accuracy

How does the AI SOC Analyst vendor ensure the AI's output is accurate, unbiased, and explainable?



— Data privacy and security —

33. Model training on customer data

Are there guarantees that customer raw data won't be used for model training without consent? What contractual commitments prevent customer data from being shared with third parties or AI model providers?



Ready to supercharge
your security operations?

REQUEST A DEMO